

# Formulating RDF Constraints Generically and Validating RDF Data Out-Of-The-Box

Thomas Bosch  
GESIS – Leibniz Institute  
for the Social Sciences, Germany  
thomas.bosch@gesis.org

Kai Eckert  
Stuttgart Media University, Germany  
eckert@hdm-stuttgart.de

## Abstract

In the context of the DCMI RDF Application Profile task group and the W3C Data Shapes Working Group solutions for the proper formulation and validation of constraints on RDF data are developed. Several approaches and constraint languages exist but there is no clear favorite and none of the languages is able to meet all requirements raised by data practitioners.

To support the work, a comprehensive, community-driven database has been created where case studies, use cases, requirements and solutions are collected. Based on this database, we published by today 81 types of constraints that are required by various stakeholders for data applications. We generally use this collection of constraint types to gain a better understanding of the expressiveness of existing solutions and gaps that still need to be filled.

Regarding the implementation of constraint languages, we already proposed to use high-level languages to describe the constraints, but map them to SPARQL queries in order to execute the actual validation; we demonstrated this approach for Description Set Profiles.

In this paper, we generalize from the experience of implementing Description Set Profiles by introducing an abstraction layer that is able to describe any constraint type in a way that is more or less straight-forwardly transformable to SPARQL queries. It provides a basic terminology and classification system for RDF constraints to lay the ground for discussions on RDF validation.

We demonstrate that using another layer on top of SPARQL helps to implement validation consistently across constraint languages and simplifies the mapping of new languages.

**Keywords:** RDF validation; RDF constraints; RDF constraint types, RDF validation requirements; Linked Data; Semantic Web

## 1 Introduction

Recently, RDF validation as a research field gained speed due to common needs of data practitioners. A typical example is the library domain that co-developed and adopted

**Thomas:**  
solutions  
for the  
proper for-  
mulation of  
constraints  
and vali-  
dation of  
RDF data  
on these  
constraints  
are devel-  
oped

**Thomas:**  
ich denke  
hier wird  
nicht klar  
was es be-  
deutet neue  
sprachen  
zu mappen,  
auch nicht  
worauf und  
warum /  
Vorschlag:  
and sim-  
plifies the  
mapping  
of new lan-

Linked Data principles very early. For libraries, the common description of resources are key business and they have a long tradition in developing and using interoperable data formats. While they embrace the openness of Linked Data and the data modeling principles provided by RDF, the data is still mostly represented in XML and this is unlikely to change any time soon. Among the reasons for the success of XML is the possibility to formulate fine-grained constraints to be met by the data and to validate the data according to these constraints using powerful systems like *DTD*, *XML Schema*, *RELAX NG*, or *Schematron*. A typical example is the definition of a library record describing a book. There are clear rules which information has to be available to describe a book properly, but also how information like an ISBN number is properly represented. Libraries seek to make their own data reusable for general purposes, but also to enrich and interlink their own data. Checking if third-party data meets own requirements or validating existing data according to new needs for a Linked Data application are among common use cases for RDF validation.

In 2013, the *W3C* organized the *RDF Validation Workshop*<sup>1</sup> where experts from industry, government, and academia discussed first RDF validation use cases. In 2014, two working groups on RDF validation have been established: the *W3C RDF Data Shapes Working Group*<sup>2</sup> and the *DCMI RDF Application Profiles Task Group*<sup>3</sup>. We collected the findings of these working groups and initiated a database of RDF validation requirements<sup>4</sup> with the intention to collaboratively collect case studies, use cases, requirements, and solutions in a comprehensive and structured way (Bosch & Eckert, 2014a). Based on our work in the *DCMI* and in cooperation with the *W3C* working group, we published by today 81 requirements to validate RDF data and to formulate constraints; each of them corresponds to a constraint type from which concrete constraints are instantiated to be checked on RDF data. We recently published a technical report (serving as first appendix of this paper) in which we explain each constraint type in detail and give examples for each represented by different constraint languages (Bosch, Nolle, Acar, & Eckert, 2015). More than 10 constraint languages with different syntax and semantics exist or are currently developed. We focus on the five most promising ones on being the standard like *Shape Expressions (ShEx)*, *Resource Shapes (ReSh)*, and *Description Set Profiles (DSP)*. The *Web Ontology Language (OWL)* in its current version *OWL 2* can also be used as a constraint language. With its direct support of validation via *SPARQL*, the *SPARQL Inferencing Notation (SPIN)* is very popular and certainly plays an important role for the future development in this field.

## 1.1 Motivation and Overview

We evaluated which constraint types the five most common constraint languages enable to express (Bosch et al., 2015). Table 1 shows in percentage values (and absolute numbers in brackets) how many constraint types each language supports.

TABLE 1: Constraint Type Specific Expressivity of Constraint Languages

<i>DSP</i>	<i>ReSh</i>	<i>ShEx</i>	<i>OWL 2</i>	<i>SPIN</i>
17.3 (14)	25.9 (21)	29.6 (24)	67.9 (55)	<b>100.0 (81)</b>

The fact that *SPIN* supports all and *OWL 2* covers 2/3 of all constraint types emphasizes the significant role *SPIN* and *OWL 2* play for the future development of constraint

<sup>1</sup><http://www.w3.org/2012/12/rdf-val/>

<sup>2</sup><http://www.w3.org/2014/rds/chapter>

<sup>3</sup><http://wiki.dublincore.org/index.php/RDF-Application-Profiles>

<sup>4</sup>Online available at <http://purl.org/net/rdf-validation>

languages. *SPARQL* and therefore *SPIN*, the *SPARQL*-based way to formulate and check constraints, is generally seen as the method of choice to validate data according to certain constraints (Fürber & Hepp, 2010), although, it is not ideal for their formulation. In contrast, high-level constraint languages are comparatively easy to understand and constraints can be formulated more concisely. Declarative languages may be placed on top of *SPARQL* and *SPIN* when using them as implementation languages.

There exists no single best solution considered as high-level intuitive constraint language which enables to express constraints in an easy and concise way and which is able to meet all requirements raised by data practitioners. Thus, the idea behind this paper is to satisfy all RDF validation requirements by representing constraints of any constraint type in a generic way using a lightweight vocabulary which consists of only a few terms. We lately published a technical report (serving as second appendix of this paper) in which we describe the vocabulary in detail (Bosch & Eckert, 2015) (Section 2). The constraint type *minimum qualified cardinality restrictions* which corresponds to the requirement *R-75* can be instantiated to ensure that *publications* must have at least one *author* which must be a *person*. As the book *The-Lord-Of-The-Rings* is a *publication*, it must have at least one *author* relationship to a *person*. The constraint is violated either (1) if the book does not have any *author* relationship, or (2) if it has an *author* which is not a *person*, or (3) if it has an *author* whose class is unknown. This constraint can either be represented generically (*generic constraint*) by *description logics (DL)*:  $\text{Publication} \sqsubseteq \geq 1 \text{ author.Person}$  or specifically (*specific constraint*) by multiple domain-specific constraint languages:

```

1 OWL 2: Publication a owl:Restriction ;
2     owl:minQualifiedCardinality 1 ;
3     owl:onProperty author ;
4     owl:onClass Person .
5
6 ShEx: Publication { author @Person{1, } }
7
8 ReSh: Publication a rs:ResourceShape ; rs:property [
9     rs:propertyDefinition author ;
10    rs:valueShape Person ;
11    rs:occurs rs:One-or-many ; ] .
12
13 DSP: [ dsp:resourceClass Publication ; dsp:statementTemplate [
14     dsp:minOccur 1 ; dsp:maxOccur "infinity" ;
15     dsp:property author ;
16     dsp:nonLiteralConstraint [ dsp:valueClass Person ] ] ] .
17
18 SPIN: CONSTRUCT { [ a spin:ConstraintViolation ... ] } WHERE {
19     ?this ?p ?o ; a ?C .
20     BIND ( qualifiedCardinality( ?this, ?p, ?C ) AS ?c ) .
21     BIND( STRDT ( STR ( ?c ), xsd:nonNegativeInteger ) AS ?cardinality ) .
22     FILTER ( ?cardinality < ?minimumCardinality ) . }
23
24 SPIN function qualifiedCardinality:
25 SELECT ( COUNT ( ?arg1 ) AS ?c ) WHERE { ?arg1 ?arg2 ?o . ?o a ?arg3 . }
```

The knowledge representation formalism *DL*, with its well-studied theoretical properties, provides the foundational basis to express constraints in a generic way. For this reason, we mapped constraint types to *DL* in order to determine which *DL* constructs are needed to express constraint types (Bosch et al., 2015).

The main contributions of this paper are: (1) We provide a basic terminology and classification system for RDF constraints to lay the ground for discussions on RDF validation (Section 2). (2) As current high-level constraint languages do not support all constraint

types, we developed a vocabulary to describe constraints of any constraint type in a generic way and specified its underlying semantics (Section 2). (3) We show how to enhance the interoperability of constraints expressed by different languages by using the proposed vocabulary to intermediately represent constraints which enables transformations between semantically equivalent constraints (Section 1.2). (4) We explain how constraint languages may reuse validation implementations which are provided once for each constraint type (Section 1.3).

## 1.2 How to Improve the Interoperability of RDF Constraints

RDF data providers have to ensure high data quality and therefore that their data conforms to constraints which may be part of contracts between data providers and consumers. As there is no standard way to formulate constraints, semantically equivalent *specific constraints*, so the *minimum qualified cardinality restriction*, may be represented by a variety of languages - each of them having different syntax and semantics. This causes confusion and weakens the common understanding between several parties about the meaning of particular constraints. We propose transformations between semantically equivalent *specific constraints* (1) to avoid the necessity to understand several languages, (2) to resolve misunderstandings about the meaning of particular constraints within the communication of RDF data producers and consumers, and (3) to enhance the interoperability of constraint languages.

We suggest to transform a *specific constraint* ( $sc_\alpha$ ) of any constraint type expressed by language  $\alpha$  into a semantically equivalent *specific constraint* ( $sc_\beta$ ) of the same constraint type represented by any other language  $\beta$  by using the developed vocabulary to intermediately represent constraints in a generic way. By defining mappings between equivalent *specific constraints* and the corresponding *generic constraint* ( $gc$ ) we are able to convert them automatically:  $gc = m_1(sc_\alpha) = m_2(sc_\beta)$ . Thereby, we do not need to define mappings between each pair of semantically equivalent *specific constraints*. Let's assume we are able to express constraints of the *minimum qualified cardinality restrictions* constraint type by 10 languages. Without an intermediate generic representation of constraints, we would have to define for each constraint type one mapping between each pair of *specific constraints* expressed by  $n$  languages - that are  $\binom{n}{2}$  mappings ( $\binom{10}{2} = 45$  mappings). With an intermediate generic representation of constraints, on the other side, we only need to define for each constraint type  $n$  mappings from  $n$  *specific constraints* to the corresponding *generic constraint* (10 mappings). For the *minimum qualified cardinality restriction*, one mapping ( $m_1(sc_{OWL2}), m_2(sc_{ShEx}), m_3(sc_{ReSh}), m_4(sc_{DSP}), \dots$ ) is defined for each *specific constraint* ( $sc_{OWL2}, sc_{ShEx}, sc_{ReSh}, sc_{DSP}, \dots$ ) to the corresponding *generic constraint* ( $gc$ ).

## 1.3 How to Provide the Validation of RDF Constraints Out-Of-The-Box

We use *SPIN* as basis to develop a validation environment<sup>5</sup> to validate RDF data according to constraints of constraint types which are expressible by arbitrary constraint languages by mapping them to *SPIN*<sup>6</sup> (Bosch & Eckert, 2014b). The SPIN engine checks for each resource if it satisfies all constraints, which are associated with its assigned classes, and generates a result RDF graph containing information about all constraint violations (Bosch & Eckert, 2014b). The next step, however, was to extend the *RDF Validator* to validate

<sup>5</sup>Online demo available at <http://purl.org/net/rdfval-demo>, source code available at: <https://github.com/boschthomas/rdf-validator>

<sup>6</sup>*SPIN* mappings available at: <https://github.com/boschthomas/rdf-validation/tree/master/SPIN>

RDF data conforming to constraints of any constraint type by representing constraints generically using the proposed vocabulary.

When language designers extend constraint languages to be able to formulate constraints of not yet supported constraint types, our proposal prevents the necessity to actually implement the validation of data according to constraints of these constraint types for each of those languages. For any constraint language, we enable to offer a validation implementation of any constraint type out-of-the-box by providing a *SPIN* mapping for each constraint type<sup>7</sup> whose constraints are represented generically (Bosch & Eckert, 2015). All what language designers need to do is to specify for each constraint type they want to cover one mapping from the *specific constraint* expressed by the language to the corresponding *generic constraint*. 19 of the overall 81 constraint types are expressible by at least four constraint languages (Bosch et al., 2015). This means that for these 19 constraint types there must be 76 implementations (four for each constraint type) to actually validate RDF data according to 76 *specific constraints* - instead of only 19 implementations (one for each *generic constraint*). Our proposal makes sure that the validation on semantically equivalent *specific constraints* expressed by different languages leads to exactly the same validation results, i.e., that validation is performed independently from the used language. This is not the case, however, for existing languages due to differences in semantics and since they are implemented independently from each other.

## 2 A Vocabulary to Describe RDF Constraints Generically

Prerequisites to develop a vocabulary to describe constraints of any constraint type in a generic way are (1) to specify the underlying semantics for validation and thus for the vocabulary and (2) to define a basic terminology and classification system for RDF constraints which lay the ground for discussions on validation.

### 2.1 Semantics for RDF Validation

RDF validation and *OWL 2* assume different semantics. This ambiguity in semantics is one of the main reasons why *OWL 2* has not been adopted as a standard constraint language for RDF validation in the past. We compare semantics of RDF validation and *OWL 2* as *OWL 2* is considered as a constraint language when validation semantics is applied. *OWL 2* is based on the *open-world assumption* (*OWA*), i.e., a statement cannot be inferred to be false if it cannot be proved to be true which fits its primary design purpose to represent knowledge on the *WWW*. As each *book* must have a *title* ( $\text{Book} \sqsubseteq \exists \text{title}.\top$ ) and *Hamlet* is a *book*, *Hamlet* must have at least one *title*. In an *OWA* setting, this constraint does not cause a violation, even if there is no explicitly defined *title*, since there must be a *title* for this *book* which we may not know. As RDF validation has its origin in the XML world many RDF validation scenarios require the *closed-world assumption* (*CWA*), i.e., a statement is inferred to be false if it cannot be proved to be true. Thus, classical constraint languages are based on the *CWA* where constraints need to be satisfied only by named individuals. The *CWA* yields to a violation since there is no explicitly defined *title* for the *book Hamlet*. Validation requires that different names represent different objects (*unique name assumption* (*UNA*)), whereas, *OWL 2* is based on the *non-unique name assumption* (*nUNA*). If we define the property *title* to be functional ( $\text{funct}(\text{title})$ ), a *book* can have at most one distinct *title*. *UNA* causes a clash if the *book Huckleberry-Finn* has more than one *title*. For *nUNA*, however, reasoning concludes that the *title The-Adventures-of-*

<sup>7</sup>RDF-CV to SPIN online available at: <https://github.com/boschthomas/RDF-CV-2-SPIN>

*Huckleberry-Finn* must be the same as the *title Die-Abenteuer-des-Huckleberry-Finn* which resolves the violation. If validation would assume *OWA* and *nUNA*, validation won't be that restrictive and therefore we won't get the intended validation results. Differences in semantics may lead to differences in validation results when applied to particular constraint types. This is important as for 56.8% of the constraint types validation results differ if the *CWA* or the *OWA* and as for 66.6% of the constraint types validation results are different in case the *UNA* or the *nUNA* is assumed (Bosch et al., 2015).

Constraints are identical w.r.t. RDF semantics, if they detect the same set of violations regardless of RDF data, which means whenever the constraints are applied to any RDF data they point out the same violations. As we use multiple languages to express constraints of constraint types, we thereby document identical semantics of these languages with regard to given constraint types (Bosch et al., 2015). This is also important to prove that semantically equivalent *specific constraints* and corresponding *generic constraints* behave in the same way w.r.t. validation results which semantically underpins bidirectional mappings.

## 2.2 Basic Terminology and Classification System for RDF Constraints

A *constraint language* is a language which is used to formulate constraints. The *W3C* working group defines *RDF constraint* as a component of a schema what needs to be satisfied<sup>8</sup>. As there are already five promising constraint languages, our purpose was not to invent a new language. As there is no high-level declarative language which meets all RDF validation requirements, we rather developed a very simple lightweight vocabulary, consisting of only three classes, three object properties, and three data properties, which is universal enough to describe constraints of any constraint type. We call this vocabulary the *RDF Constraints Vocabulary (RDF-CV)*<sup>9</sup> whose conceptual model is shown in Figure 1.




FIG. 1: *RDF Constraints Vocabulary (RDF-CV)* Conceptual Model

We classify sets of constraints according to four dimensions. For each constraint type, we evaluated (Bosch et al., 2015) in which set of constraints their instances are contained. **Universality:** Constraints are expressible either generically (*generic constraints*) using the *RDF-CV* or specifically (*specific constraints*) by domain-specific constraint languages. As we express constraints of each constraint type in form of *generic constraints* conforming to the *RDF-CV* (Bosch et al., 2015), we show that constraints of any constraint type can be described generically. **DL Expressivity:** The *RDF-CV* allows to describe constraints which are *expressible in DL* (constraints of this set are instantiated from 64% of the overall 81 constraint types) and which are *not expressible in DL* (36%). **Complexity:** *Simple constraints* (60% of the constraint types) denotes the set of atomic constraints. *Complex constraints* (26%) is the set of constraints which are created out of *simple* and/or other *com-*

<sup>8</sup><https://www.w3.org/2014/data-shapes/wiki/Glossary>

<sup>9</sup>Online available at: <https://github.com/boschthomas/RDF-Constraints-Vocabulary>

*plex constraints*. Constraints of almost 14% of the constraint types are *complex constraints* which can be simplified and therefore formulated as *simple constraints* when using them in terms of syntactic sugar. **Context:** *Simple constraints* are applied on properties (*property constraints*, 60%), on classes (*class constraints*, 25%), or on properties and classes (*property and class constraints*, 15%). *Simple constraints* must hold for individuals of their *context classes*. As *context classes* of *simple constraints* are reused within *complex constraints*, it does not make sense to have terms standing for *simple* and *complex constraints* within the *RDF-CV*. As a consequence, the distinction of *property* and *class constraints* is sufficient to generically describe constraints of all constraint types. Altogether, instances of the most of the constraint types are *simple constraints on properties* which are *expressible in DL*.

## 2.3 Simple Constraints

The implementation model of the *RDF-CV* (Figure 2) shows that both *property constraints* and *class constraints* are *simple constraints*.




FIG. 2: *RDF Constraints Vocabulary (RDF-CV)* Implementation Model

A *simple constraint* holds for all individuals of their *context class*. The *minimum qualified cardinality restriction*  $\text{Publication} \sqsubseteq \geq 1 \text{ author.Person}$ , which restricts *publications* to have at least one *author* which must be a *person*, is an example of a *property constraint* on *author* which holds for all individuals of the class *Publication*. Table 2 displays the mapping to the *RDF-CV*.

TABLE 2: Minimum Qualified Cardinality Restriction as Property Constraint

constraint set	context class	left property list	right p. list	classes	constraining element	c. value
property	Publication	author	-	Person	$\geq$	1

The *constraining element* indicates the actual type of constraint like *DL* concept and role constructors, (in)equality, and further keywords for constraint types whose constraints cannot be expressed in *DL* (e.g., regular expressions and constraining facets). In some cases, a constraint is only complete when a *constraining value* is stated in addition to the *constraining element* and *simple constraints* may refer to a list of *classes*. The *constraining element* of the *property constraint*  $\text{Publication} \sqsubseteq \geq 1 \text{ author.Person}$ , e.g., is the *DL at-least restriction*  $\geq$ , the *constraining value* is 1, and the list of *classes* includes *Person* which restricts the objects of the property *author* to be of the class *Person*. For *property constraints*, *left* and *right property lists* are specified. The assignment of properties to these lists happens relative to the *constraining element* which may be an operator like  $\sqsubseteq$ . *Object Property Paths* (R-55) ensure that if an individual  $x$  is connected by a sequence of object properties with an individual  $y$ , then  $x$  is also related to  $y$  by a particular object property. As *Stephen-Hawking* is the *author* of the book *A-Brief-History-Of-Time* whose *genre* is *Popular-Science*, the *object property path*  $\text{authorOf} \circ \text{genre} \sqsubseteq \text{authorOfGenre}$  infers that *Stephen-Hawking* is an *author* of the *genre Popular-Science*. Thus, when mapped to the *RDF-CV* (Table 3), the properties *authorOf* and *genre* are placed on the left side of the

*constraining element*  $\sqsubseteq$  and the property *authorOfGenre* on the right side. As this *property constraint* holds for all individuals within the data, the *context class* is the *DL top concept*  $\top$  which stands for the super-class of all classes.

TABLE 3: Object Property Paths as Property Constraint

c. set	context class	left p. list	right p. list	classes	c. element	c. value
property	$\top$	authorOf, genre	authorOfGenre	$\top$	$\sqsubseteq$	-

There are *simple constraints* which are not expressible in *DL* but can still be described using the *RDF-CV* such as constraints of the type *literal pattern matching* (R-44) which restrict literals to match given patterns. The *universal restriction*  $\text{Book} \sqsubseteq \forall \text{identifier.ISBN}$  ensures that *books* can only have valid *ISBN* identifiers, i.e., strings that match a given regular expression. Table 4 presents how the *literal pattern matching simple constraint* which is not expressible in *DL* and how the *universal restriction complex constraint* which is expressible in *DL* are mapped to the *RDF-CV*. Thereby, the *context class* *ISBN*, whose instances must satisfy the *simple constraint*, is reused within the list of *classes* the *complex constraint* refers to. The *literal pattern matching* constraint type introduces the *constraining element* *regex* whose validation has to be implemented once like for any other *constraining element*.

TABLE 4: Simple Constraints which are not Expressible in DL)

c. set	context class	left p. list	right p. list	classes	c. element	c. value
class	ISBN	-	-	xsd:string	regex	'^d{9}[d X]\$',
property	Book	identifier	-	ISBN	$\forall$	-

## 2.4 Complex Constraints

*Complex constraints* of the constraint type *context-specific exclusive or of property groups* (R-13) restrict individuals of given classes to have all properties of exactly one of multiple mutually exclusive property groups. *Publications*, e.g., are either identified by an *ISBN* and a *title* (for books) or by an *ISSN* and a *title* (for periodical *publications*), but it should not be possible to assign both identifiers to a given *publication*. As *The-Great-Gatsby* is a *publication* with an *ISBN* and a *title* without an *ISSN*, *The-Great-Gatsby* is considered as a valid *publication*. The *complex constraint* is generically expressible by *DL*:

$$\begin{aligned} \text{Publication} &\sqsubseteq (\neg \mathbf{E} \sqcap \mathbf{F}) \sqcup (\mathbf{E} \sqcap \neg \mathbf{F}), \mathbf{E} \equiv \mathbf{A} \sqcap \mathbf{B}, \mathbf{F} \equiv \mathbf{C} \sqcap \mathbf{D} \\ \mathbf{A} &\sqsubseteq \geq 1 \text{ isbn.string} \sqcap \leq 1 \text{ isbn.string}, \mathbf{B} \sqsubseteq \geq 1 \text{ title.string} \sqcap \leq 1 \text{ title.string} \\ \mathbf{C} &\sqsubseteq \geq 1 \text{ issn.string} \sqcap \leq 1 \text{ issn.string}, \mathbf{D} \sqsubseteq \geq 1 \text{ title.string} \sqcap \leq 1 \text{ title.string} \end{aligned}$$

The *DL* statements demonstrate that the *complex constraint* is composed of many other *complex* (e.g. *minimum and maximum qualified cardinality restrictions*) and *simple constraints* (e.g. constraints on sets). As *exact (un)qualified cardinality restrictions* ( $=n$ ) and *exclusive or (XOR)* are frequently used *complex constraints*, we propose to simplify them as *simple constraints* in terms of syntactic sugar. As a consequence, the *context-specific exclusive or of property groups complex constraint* is represented as a *generic constraint* by means of the *RDF-CV* more intuitively and concisely (Table 5).



TABLE 5: Simplified Complex Constraints

c. set	context class	left p. list	right p. list	classes	c. element	c. value
class	Publication	-	-	E, F	XOR	
class	E	-	-	A, B	$\sqcap$	
class	F	-	-	C, D	$\sqcap$	
property	A	isbn	-	string	=	1
property	B	title	-	string	=	1
property	C	issn	-	string	=	1
property	D	title	-	string	=	1

The *primary key properties* constraint type (*R-226*) is often useful to declare a given (datatype) property as the primary key of a class, so that a system can enforce uniqueness. *Books*, e.g., are uniquely identified by their *ISBN*, i.e., the property *isbn* is inverse functional (`funct isbn-`) which can be represented using the *RDF-CV* in form of a *complex constraint* consisting of two *simple constraints* (Table 6).

TABLE 6: Primary Key Properties as Complex Constraints

c. set	context class	left p. list	right p. list	classes	c. element	c. value
property	$\top$	isbn <sup>-</sup>	isbn <sup>-</sup>	-	inverse	-
property	Book	isbn <sup>-</sup>	-	-	functional	-

*Keys*, however, are even more general, i.e., a generalization of *inverse functional properties* (Schneider, 2009). A key can be a datatype, an object property, or a chain of properties. For these generalization purposes, as there are different sorts of keys, and as keys can lead to undecidability, *DL* is extended with a special construct *keyfor* (Lutz, Areces, Horrocks, & Sattler, 2005). When using *keyfor* (`isbn keyfor Book`), the *complex constraint* can be formulated as a *simple constraint* which looks like the following in concrete RDF turtle syntax:

```

1 [ a rdfcv:PropertyConstraint , rdfcv:SimpleConstraint ;
2   rdfcv:contextClass Book ; rdfcv:leftProperties ( isbn ) ; rdfcv:constrainingElement "keyfor" ] .

```

*Complex constraints* of frequently used constraint types corresponding to *DL* axioms like *transitivity*, *symmetry*, *asymmetry*, *reflexivity* and *irreflexivity* can also be simplified in the same way, as, although, such axioms can be expressed using basic *DL* features, *DL* axioms like *transitivity* can also be used in terms of syntactic sugar (Bosch & Eckert, 2015).

### 3 Related Work

In this section, we present current languages for constraint formulation and RDF data validation<sup>10</sup>. *SPIN*, *SPARQL*, *OWL 2*, *ShEx*, *ReSh*, and *DSP* are the six most promising and applied languages. The query language for RDF *SPARQL* is generally seen as the method of choice to validate data according to certain constraints (Fürber & Hepp, 2010), although, it is not ideal for their formulation. In contrast, high-level constraint languages are comparatively easy to understand and constraints can be formulated more concisely. Declarative languages may be placed on top of *SPARQL* and *SPIN* when using them as

<sup>10</sup>*SPARQL*: <http://www.w3.org/TR/sparql11-query>, *SPIN*: <http://spinrdf.org>, *Pellet ICV*: <http://clarkparsia.com/pellet/icv>, *Stardog ICV*: [http://docs.stardog.com/#\\_validating\\_constraints](http://docs.stardog.com/#_validating_constraints), *ShEx*: <http://www.w3.org/Submission/shex-defn>, *ReSh*: <http://www.w3.org/Submission/shapes>, *DCMI RDF-AP*: <http://dublincore.org/documents/singapore-framework>, *Bibframe*: <http://bibframe.org>, *DSP*: <http://dublincore.org/documents/dc-dsp>, *RDF Schemarama*: <http://swordfish.rdfweb.org/discovery/2001/01/schemarama>

implementation languages. *SPIN* provides a vocabulary to represent *SPARQL* queries as RDF triples and uses *SPARQL* to specify logical constraints and inference rules (Fürber & Hepp, 2010). *Stardog ICV* and *Pellet ICV* use *OWL 2* constructs to formulate constraints. The *Pellet Integrity Constraint Validator (ICV)* is a proof-of-concept extension for the *OWL 2 DL* reasoner *Pellet*. *Stardog ICV* validates RDF data stored in a *Stardog* database according to constraints which may be written in *SPARQL*, *OWL*, or *SWRL*. *ShEx* specifies a language whose syntax and semantics are similar to regular expressions. *ShEx* associate RDF graphs with labeled patterns called *shapes* which are used to express formal constraints on the content of RDF graphs. *ReSh* defines its own vocabulary for specifying *shapes* of RDF resources. Ryman et al. define *shape* as a description of the set of triples a resource is expected to contain and of the integrity constraints those triples are required to satisfy (Ryman, Hors, & Speicher, 2013). *DCMI RDF Application Profiles (AP)* and *Bibframe* are approaches to specify profiles for application-specific purposes. *DCMI RDF-AP* uses *DSP* as generic constraint language which is also intuitive for non-experts. The *Bibliographic Framework Initiative (Bibframe)* defines a vocabulary which has a strong overlap with *DSP*. Kontokostas et al. define 17 data quality integrity constraints represented as *SPARQL* query templates called *Data Quality Test Patterns (DQTP)* (Kontokostas et al., 2014). *RDF Schemarama* is a proof-of-concept demonstrator based on ideas from *Schematron*. It allows to check that RDF data has required properties and it is based on *Squish*, an SQL-ish query language for RDF, instead of *SPARQL*. In addition to the formulation of constraints, *SPIN* (open source API), *Stardog ICV* (as part of the *Stardog RDF database*), *DQTP* (tests), *Pellet ICV* (extension of *Pellet OWL 2 DL* reasoner) and *ShEx* offer executable validation systems using *SPARQL* as implementation language.

## 4 Conclusion and Future Work

Based on our work in the *DCMI* and in cooperation with the *W3C* working group, we published by today 81 requirements to validate RDF data and to formulate constraints; each of them corresponds to a constraint type from which concrete constraints are instantiated to be checked on RDF data. These constraint types form the basis to lay the ground for discussions on validation by defining a basic terminology and classification system for RDF constraints. There exists no single best solution considered as high-level intuitive constraint language which enables to express constraints in an easy and concise way and which is able to meet all requirements raised by data practitioners. Thus, the idea behind this paper is to satisfy all requirements by representing constraints of any constraint type in a generic way using a lightweight vocabulary which consists of only a few terms.

As there is no standard way to formulate constraints, semantically equivalent *specific constraints* may be represented by a variety of languages - each of them having different syntax and semantics. We propose transformations between semantically equivalent *specific constraints* by using the proposed vocabulary to intermediately represent constraints in a generic way (1) to avoid the necessity to understand several languages, (2) to resolve misunderstandings about the meaning of particular constraints, and (3) to enhance the interoperability of constraint languages (Section 1.2).

We use *SPIN* as basis to develop a validation environment<sup>5</sup> to validate RDF data according to constraints of constraint types which are expressible by arbitrary constraint languages by mapping them to *SPIN*<sup>6</sup> (Bosch & Eckert, 2014b). When language designers extend constraint languages to be able to formulate constraints of not yet supported constraint types, our proposal enables to reuse the validation implementation of these constraint types by mapping *specific constraints* expressed by the language to the corresponding

*generic constraint*. For any constraint language, we enable to offer a validation implementation of any constraint type out-of-the-box by providing a *SPIN* mapping for each constraint type<sup>7</sup> whose constraints are represented generically (Bosch & Eckert, 2015). Furthermore, our proposal makes sure that the validation on semantically equivalent *specific constraints* expressed by different languages leads to exactly the same validation results, i.e., that validation is performed independently from the used language (Section 1.3).

It is part of future work (1) to extend the *RDF Validator* by generating *generic constraints* according to inputs of domain experts who may not be familiar with the formulation of constraints, (2) to offer bidirectional transformations between *specific constraints* expressed by the most common constraint languages and corresponding *generic constraints*, and (3) to provide translations between semantically equivalent *specific constraints* expressed by the most common constraint languages by using the developed vocabulary to intermediately represent constraints. Identifying RDF validation requirements is an ongoing process, so the task to map constraints of new constraint types to the vocabulary.

## References

- Bosch, T., & Eckert, K. (2014a). Requirements on rdf constraint formulation and validation. *Proceedings of the DCMi International Conference on Dublin Core and Metadata Applications (DC 2014)*.
- Bosch, T., & Eckert, K. (2014b). Towards description set profiles for rdf using sparql as intermediate language. *Proceedings of the DCMi International Conference on Dublin Core and Metadata Applications (DC 2014)*.
- Bosch, T., & Eckert, K. (2015). A vocabulary to describe rdf constraints generically.
- Bosch, T., Nolle, A., Acar, E., & Eckert, K. (2015). Rdf validation requirements - evaluation and logical underpinning.
- Fürber, C., & Hepp, M. (2010). Using sparql and spin for data quality management on the semantic web. In W. Abramowicz & R. Tolksdorf (Eds.), *Business information systems* (Vol. 47, pp. 35–46). Springer Berlin Heidelberg. Retrieved from [http://dx.doi.org/10.1007/978-3-642-12814-1\\_4](http://dx.doi.org/10.1007/978-3-642-12814-1_4) doi: 10.1007/978-3-642-12814-1\_4
- Kontokostas, D., Westphal, P., Auer, S., Hellmann, S., Lehmann, J., Cornelissen, R., & Zaveri, A. (2014). Test-driven evaluation of linked data quality. In *Proceedings of the 23rd international conference on world wide web* (pp. 747–758). Republic and Canton of Geneva, Switzerland: International World Wide Web Conferences Steering Committee. Retrieved from <http://dx.doi.org/10.1145/2566486.2568002> doi: 10.1145/2566486.2568002
- Lutz, C., Areces, C., Horrocks, I., & Sattler, U. (2005, June). Keys, nominals, and concrete domains. *Journal of Artificial Intelligence Research*, 23(1), 667–726. Retrieved from <http://dl.acm.org/citation.cfm?id=1622503.1622518>
- Ryman, A. G., Hors, A. L., & Speicher, S. (2013). Oslc resource shape: A language for defining constraints on linked data. In C. Bizer, T. Heath, T. Berners-Lee, M. Hausenblas, & S. Auer (Eds.), *Ldow* (Vol. 996). CEUR-WS.org. Retrieved from <http://dblp.uni-trier.de/db/conf/www/ldow2013.html#RymanHS13>
- Schneider, M. (2009, October). *OWL 2 Web Ontology Language RDF-Based Semantics* (W3C Recommendation). W3C. Retrieved from [\url{http://www.w3.org/TR/2009/REC-owl2-rdf-based-semantics-20091027/}](http://www.w3.org/TR/2009/REC-owl2-rdf-based-semantics-20091027/)